

# 1. Importing necessary libraries

```
In [1]: import pandas as pd  
import numpy as np  
from matplotlib import pyplot as plt  
import seaborn as sns
```

## 2. Importing data



```
In [2]: salary_data = pd.read_csv("Salary_Data.csv")
salary_data
```

```
Out[2]:
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

### 3. Data Understanding

```
In [3]: salary_data.shape
```

```
Out[3]: (30, 2)
```

```
In [4]: salary_data.isna().sum()
```

```
Out[4]: YearsExperience    0  
Salary                    0  
dtype: int64
```

```
In [5]: salary_data.dtypes
```

```
Out[5]: YearsExperience    float64  
Salary                    float64  
dtype: object
```

```
In [6]: salary_data.describe(include='all')
```

```
Out[6]:
```

	YearsExperience	Salary
<b>count</b>	30.000000	30.000000
<b>mean</b>	5.313333	76003.000000
<b>std</b>	2.837888	27414.429785
<b>min</b>	1.100000	37731.000000
<b>25%</b>	3.200000	56720.750000
<b>50%</b>	4.700000	65237.000000
<b>75%</b>	7.700000	100544.750000
<b>max</b>	10.500000	122391.000000

## 4. Check Assumptions are matching

### 1. Check for linearity.

```
In [7]: sns.scatterplot(x = 'YearsExperience',y= 'Salary',data=salary_data)
plt.title('YearsExperience VS Salary',size = 17)
plt.show()
```



```
In [8]: sns.lmplot(x = 'YearsExperience',y= 'Salary',data=salary_data)
plt.title('YearsExperience VS Salary',size = 17)
plt.show()
```



**Observation = Linearity check is passed**

```
In [9]: salary_data.corr().round(2)
```

```
Out[9]:
```

	YearsExperience	Salary
YearsExperience	1.00	0.98
Salary	0.98	1.00

**2.Homoscadasticity-it can be checked post model building and training**

**3.NO multicollinarity**

**4. No Autoregression**

**5. zero residual mean-it can be checked post model building and training**

**5.Data preparation**

```
In [10]: salary_data #no unwanted parameters.
```

```
Out[10]:
```

	<b>YearsExperience</b>	<b>Salary</b>
<b>0</b>	1.1	39343.0
<b>1</b>	1.3	46205.0
<b>2</b>	1.5	37731.0
<b>3</b>	2.0	43525.0
<b>4</b>	2.2	39891.0
<b>5</b>	2.9	56642.0
<b>6</b>	3.0	60150.0
<b>7</b>	3.2	54445.0
<b>8</b>	3.2	64445.0
<b>9</b>	3.7	57189.0
<b>10</b>	3.9	63218.0
<b>11</b>	4.0	55794.0
<b>12</b>	4.0	56957.0
<b>13</b>	4.1	57081.0
<b>14</b>	4.5	61111.0
<b>15</b>	4.9	67938.0
<b>16</b>	5.1	66029.0
<b>17</b>	5.3	83088.0
<b>18</b>	5.9	81363.0
<b>19</b>	6.0	93940.0
<b>20</b>	6.8	91738.0
<b>21</b>	7.1	98273.0
<b>22</b>	7.9	101302.0
<b>23</b>	8.2	113812.0
<b>24</b>	8.7	109431.0
<b>25</b>	9.0	105582.0
<b>26</b>	9.5	116969.0
<b>27</b>	9.6	112635.0
<b>28</b>	10.3	122391.0
<b>29</b>	10.5	121872.0

## 6.Model Building and Model training

There are basically 2 libraries that supports Linear Regression algorithm.

- 1.Statsmodels libraries-ols techniques.
- 2.sklearn libraries-linear regression.

```
In [11]: import statsmodels.formula.api as smf
```

```
In [21]: lin_reg_model = smf.ols(formula='Salary~YearsExperience', data=salary_data).fit()
```

note- with stats models model building and training is happening at the same time

## 7.Model Testing

```
In [22]: lin_reg_model.params
```

```
Out[22]: Intercept      25792.200199  
YearsExperience    9449.962321  
dtype: float64
```

### Manual Testing

```
In [15]: # y = mx + c  
#for x = 1.3 y=?
```

```
In [23]: 9449.962321 *1.1 +25792.200199
```

```
Out[23]: 36187.158752100004
```

```
In [25]: 9449.962321 *1.3 +25792.200199
```

```
Out[25]: 38077.1512163
```

```
In [26]: 9449.962321 *1.5 +25792.200199
```

```
Out[26]: 39967.1436805
```

```
In [27]: 9449.962321 *2 +25792.200199
```

```
Out[27]: 44692.124841
```

```
In [28]: 9449.962321 *2.2 +25792.200199
```

```
Out[28]: 46582.1173052
```

### Machine Prediction



```
In [30]: test_data = pd.read_csv('salary_test_data.csv')
test_data
```

```
Out[30]:
```

	YearsExperience
0	1.1
1	1.3
2	1.5
3	2.0
4	2.2

```
In [31]: lin_reg_model.predict(test_data)
```

```
Out[31]: 0    36187.158752
1    38077.151217
2    39967.143681
3    44692.124842
4    46582.117306
dtype: float64
```

## 8. Model Evaluation

```
In [32]: lin_reg_model.rsquared, lin_reg_model.rsquared_adj
```

```
Out[32]: (0.9569566641435086, 0.9554194021486339)
```

```
In [33]: lin_reg_model.aic, lin_reg_model.bic ##is an estimator of prediction error
```

```
Out[33]: (606.882316930432, 609.6847116937563)
```

## 9. Model Deployment

```
In [34]: from pickle import dump
```

```
In [35]: dump(lin_reg_model, open('linear_regression.pkl', 'wb'))
```

```
In [36]: from pickle import load
```

```
In [37]: tested_linear_model = load(open('linear_regression.pkl', 'rb'))
```

```
In [38]: tested_linear_model.predict(test_data)
```

```
Out[38]: 0    36187.158752  
         1    38077.151217  
         2    39967.143681  
         3    44692.124842  
         4    46582.117306  
         dtype: float64
```

## Assumption check

- Homoscedaticity

```
In [39]: from sklearn.preprocessing import StandardScaler  
std_scaler = StandardScaler()  
scaled_x = std_scaler.fit_transform(salary_data)  
print(scaled_x)
```

```
[[-1.51005294 -1.36011263]  
 [-1.43837321 -1.10552744]  
 [-1.36669348 -1.419919  ]  
 [-1.18749416 -1.20495739]  
 [-1.11581443 -1.33978143]  
 [-0.86493538 -0.71830716]  
 [-0.82909552 -0.58815781]  
 [-0.75741579 -0.79981746]  
 [-0.75741579 -0.42881019]  
 [-0.57821647 -0.69801306]  
 [-0.50653674 -0.47433279]  
 [-0.47069688 -0.74976858]  
 [-0.47069688 -0.70662043]  
 [-0.43485702 -0.70201994]  
 [-0.29149756 -0.55250402]  
 [-0.1481381  -0.29921736]  
 [-0.07645838 -0.37004264]  
 [-0.00477865  0.26285865]  
 [ 0.21026054  0.19885989]  
 [ 0.2461004   0.66547573]  
 [ 0.53281931  0.58377993]  
 [ 0.6403389   0.82623317]  
 [ 0.92705781  0.93861127]  
 [ 1.03457741  1.40274136]  
 [ 1.21377673  1.24020308]  
 [ 1.32129632  1.09740238]  
 [ 1.50049564  1.51986835]  
 [ 1.5363355   1.3590738  ]  
 [ 1.78721455  1.72102849]  
 [ 1.85889428  1.70177321]]
```

```
In [40]: scaled_x.mean(),scaled_x.std()
```

```
Out[40]: (-2.2204460492503132e-17, 1.0)
```

## Zero residual mean

```
In [44]: np.mean(scaled_x)
```

```
Out[44]: -2.2204460492503132e-17
```